

Efficient Thinking VIII-c: Where the Experience Is Acquired

A frozen agent's plateau, another agent's places, a forced look, and what the difference is made of

STATE 2026-09-28: v1.0.12, 2026-09-28 (v1.0 plus a plain-words summary, a reference list, a named open question on the forced-look mechanism and a note on the gen0 arm, its wording checked against Appendix A; one second-machine reproduction recorded under Reproducibility and the first lock's head files hashed; §6a related work, all thirteen citations content-checked — six against method sections, seven at abstract level — with four characterisations corrected; withdrawn-run provenance stated precisely; no number changed). Content frozen at lock v3; every number second-read. The artefact read matched all 37 checkable numbers and the v3 delta to their artefacts; the results reader hashed every game file of locks v2 and v3 (162 of 162 and 189 of 189, each against its own lock entry) and recomputed every green count from the raw episodes (39 of 39 arm-games), re-deriving the six- and seven-arm tables from files; the hygiene read found the needles clean. Nothing in the chain from raw episodes to the tables rests on a reported number. Versions before this are drafts and are superseded.

Louay Alsakka · September 28, 2026 · v1.0.12 · doi:[10.5281/zenodo.23005271](https://doi.org/10.5281/zenodo.23005271)

In plain words. The plateau in VIII-b was not because the model had used up what it can learn. It was because of where its experience came from. Experience recorded from where another agent stood, or from a forced first look before guessing, restored gains as large as VIII first reported, with the same model and the same helper. Where you learn from matters, not just how much you collect.

Abstract

Paper VIII showed that a frozen model with a read-only head fitted on its own verified history moves the quality–compute frontier. Paper VIII-b showed that the gain does not compound: a second generation of the agent's own experience is largely a repeat of the first, and a head fitted on it is not separable from the first by the loop's own noise. This paper asks what the plateau is made of, and finds that it can arise from how the agent acquires experience rather than from any exhaustion of its ability to use it: different acquisition policies produce differently valuable experience, and acquisition is an optimisation dimension separable from the intelligence and from the mechanism that consumes experience. The evidence is three heads of identical form and size, fitted on the same frozen model's hidden states with the same verifier labels, differing only in *which states* they were fitted on. On 300 problems no head had seen: a head fitted where the

agent’s own policy stood solves 26.8% (base 22.0%); a head fitted where a second, weaker agent’s policy stood solves 35.4%; a head fitted where the agent was made to stand — the code read before any hypothesis, by a one-line rule during collection only — solves 41.2%. The second agent was the experiment that revealed the dissociation, not its mechanism: the effect needs neither another agent nor several. A head built from states the agent was made to visit is more selective than one built from states it chose — it agrees with the agent’s own pick less often, and is right more often when it does. Mixing the two acquired sets at a matched size does not beat the forced look alone, but the second agent’s groups are not dead weight: half the forced groups plus the second agent’s groups recover most of the gain, and 132 forced-look groups are indistinguishable from 290 of the second agent’s — consistent with about twice the value per group, not a measured ratio. A cheap held-out proxy tracks how much experience a head was fitted on and not where it came from, and ranked heads of equal size exactly backwards. One task family, one model pair: a mechanism demonstrated in this system, not a law.

Results at a glance

Observations — seven arms, three games on the same 300 problems (seed 89), percent of problems solved. The first four arms ran adjacent within each game in one session (09-25); the forced-look and mixture arms ran on 09-26 against that bracket (§7). Means and within-arm ranges; no comparison is implied by adjacency.

arm	what sits beside the frozen 7B	g1	g2	g3	mean	range
base	nothing	66	67	65	22.0	0.67
gen0	the head of VIII (its earliest experience)	58	60	60	19.8	0.67
self	the matched second- generation head of VIII-b (own experience)	77	83	81	26.8	2.00
other- agent	head on the 7B’s states at the places agent B’s policy visited	102	104	113	35.4	3.67
forced- look	head on the 7B’s states at the places a forced first	122	126	123	41.2	1.33

arm	what sits beside the frozen 7B	g1	g2	g3	mean	range
mixture	inspection visited					
	head on other-agent ∪ forced- look decisions, same size (drawn 158 + 132)	117	113	117	38.6	1.33
forced- look, half	head on the mixture's own 132 forced-look groups alone	95	101	97	32.6	2.00

On gen0, the only arm below base. Its head was fitted on states collected while the harness prompt carried a region menu from an earlier task family (Appendix A, step 3); on this family's corrected menu its first picks name the true region less often than the base's own (17.0% against 23.3%, §4). It is carried here as the origin of the accumulation series, not as a re-measurement of VIII: on VIII's own sets the same head scores +9 to +10 points over base (VIII-b §7) and VIII's claim is measured and reproduced in its own harness at matched compute.

Supported comparisons — paired per problem, three games, read against the registered bars (§2). Only these rows are claims.

comparison	mean difference (points)	§14(a): 95% interval excludes zero (games of 3)	§14(b): lower bound clears the session's largest within- head-arm range (games of 3)	where
other-agent – self	+8.7	3 of 3	1 of 3	§3
other-agent – base	+13.4	3 of 3	3 of 3	§3
forced-look – self	+14.4	3 of 3	3 of 3	§3
forced-look – base	+19.2	3 of 3	3 of 3	§3
self – base	+4.8	0 of 3	0 of 3	§3 (VIII-b reproduced)
mixture – forced-look	–2.7	0 of 3	0 of 3	§5 — within noise

comparison	mean difference (points)	§14(a): 95% interval excludes zero (games of 3)	§14(b): lower bound clears the session's largest within- head-arm range (games of 3)	where
forced-look, half – base	+10.6	3 of 3	3 of 3	§5
forced-look – forced-look, half	+8.7	3 of 3	0 of 3	§5 — no saturation at half size
mixture – forced-look, half	+6.0	2 of 3	0 of 3	§5 — the second agent's groups contribute
forced-look – other-agent	+5.8	2 of 3	0 of 3	§5 — within noise; not a ranking

Two rows a reader will be tempted to read as rankings are marked as within noise and are not claims: forced-look against other-agent, and mixture against forced-look. No comparison between the two acquisition sources is established here (§5).

1. The question

VIII-b ended on a diagnosis: a steered agent stands where the previous agent stood and produces the same state, so its experience stops being new. That leaves two readings of the plateau. Either the model has taken what it can from experience of this kind — its ability to *use* experience is exhausted at this size — or the model has stopped *acquiring* experience worth using, because its own policy keeps it in the same region. The two readings predict different things. If use is exhausted, no experience of this form helps further, whoever gathered it. If acquisition is the limit, experience gathered somewhere else, by any means, should help, and the head form need not change at all.

The experiment that separates them is to hold everything fixed — the frozen model, the head's form and size, the verifier's labels, the scoring set — and vary only where the states come from.

2. Instruments, and what is fixed

Everything runs on the loop harness of VIII and VIII-b (et8b_loop, budget 12, up to three head decisions per episode) with the same frozen 7B model. Three things are fixed across every arm and are the reason the comparison is clean:

- **The head.** A read-only linear head over the frozen model's hidden state at layer 18, shared across decisions, fitted on 290 whole decision groups drawn at one seed. Every head in this paper has this form and this count. Its labels are the verifier's ground truth (bug_region) in every arm; no agent supplies a label.

- **The states.** Every head is fitted on the 7B's hidden states. When another agent's trajectory is used, that trajectory is replayed as a prompt to the 7B and the 7B's own state at each of the other agent's decision points is recorded. Nothing the other agent thought enters the head; only where it stood.
- **The scoring set.** All games run on a 300-problem set (seed 89) that no head in the comparison was fitted on. The fitting sets (seed 73 for the other-agent and forced-look heads; seed 21 for the self head) share no program with it and, of 70 symptom strings the sets share, at most five that point to the same region (one of 71 for the self head's set), against effects of thirty-plus problems.

The bars are those registered in VIII-b (§14): a comparison clears §14(a) in a game when the paired bootstrap 95% interval of the per-problem difference excludes zero; it clears §14(b) when that interval's lower bound exceeds the largest within-head-arm range of the session. A row is a claim only where the table says how many of the three games clear each. The session's largest range is the other-agent arm's, 3.67 points, and it did not move when the two later arms were added (they are the steadiest arms in the table).

Two properties of the instrument, established in VIII-b and confirmed here, frame every number: decision 1 is greedy on an identical state, so a head's first pick is deterministic and identical across games; all between-game variation lives at decisions 2 and 3.

3. The dissociation

The self head reproduces the plateau. On the new set, the matched second-generation head of VIII-b beats base by +3.7, +5.3 and +5.3 points and clears no bar. This is VIII-b's result on a third problem set, and it is the control the rest of the paper is read against.

The other-agent head clears it. Agent B is a second frozen model of a different family (Llama-3.1-8B), the weaker of the two: across two collections it solves 34–36 of the 300 fitting problems where the 7B, across two runs, solves 60–68 (the head was fitted on the 36 run). Its policy differs in one measurable habit: it inspects the code before its first hypothesis on 70% of problems (209 of 300 on the run that fed the head), where the 7B rarely does. A head fitted on the 7B's states at B's visited places, labelled by the verifier, beats base by +12.0, +12.3 and +16.0 and clears both bars in every game; it beats the self head by +8.3, +7.0 and +10.7, clearing the point-estimate bar in every game and the interval bar in one. Ability to use experience was not exhausted: the same head form, on states from elsewhere, more than doubled the gain.

The other agent was not needed. The forced-look head is fitted on the 7B's own trajectories with one change during collection only: at decision 1 the action is forced to be an inspection (charged as an action, target chosen by the loop's existing rule), with decisions 2 and 3 under the native policy. That head beats base by +18.7, +19.7 and +19.3 and the self head by +15.0, +14.3 and +14.0, clearing both bars in every game. Against the other-agent head it leads in every game but the difference is within the session's noise (§5). The registered reading is that the forced look is at least as good as the second agent's own trajectories: what mattered was where the model was standing when the experience was recorded, and a one-line rule can put it there.

So the plateau of VIII-b was acquisition. The agent's own policy returned to the same region; a head fitted there was fitted where the evidence was not yet in the prompt.

4. What the difference is made of

Three pre-registered predictions were read from the games' own logs, and one alternative explanation was tested after the fact.

The head is a better function on the identical input. At decision 1 every arm sees the same state, so the share of first picks that name the true region is the head's function alone: base 23.3%, gen0 17.0%, self 30.7%, other-agent 42.3%. Of the other-agent head's 170 solves the self head did not reach, it had named the true region at decision 1 on 105 (62%); the self head had on 6 (4%). Its edge grows through decisions 2 and 3 (+11.7, +19.1, +21.3 points over the self head), but the later decisions are the dirtier comparison — by decision 2 the arms are in states their own first pick produced — so the clean number is the smallest one.

Episodes get shorter, not longer. The registered prediction that a better head would lengthen episodes (an inspection before the hypothesis) was refuted: mean actions fall, 7.59 under the other-agent head against 8.21 base, because a correct first pick ends the episode sooner. The habit itself does not transfer at game time — the loop forces a hypothesis at decision 1 under every head — and the result does not need it to.

The forced-look head is not degenerate; it is selective. The worry was that forcing every first move produces a head that always names this family's favourite region. Measured over every decision at which each head was consulted, the forced-look head is the *flattest* of the six (top-pick share 17.6% against a uniform 12.5%; the most concentrated head, gen0, is the only one that loses to base). What distinguishes it is selectivity: it agrees with the agent's own pick less often than the other-agent head (33.7% against 40.3%) and is right more often when it agrees (73.3% against 60.6%). A head built from states the agent was made to visit disagrees with the policy that produced it more, and its agreement carries more information.

5. Mixing the two, and what the matched size hides

The registered second arm fits one head on the union of the other-agent and forced-look decisions at the same size as either alone; the seed drew 158 of the second agent's groups and 132 forced-look groups. It solves 38.6%: below the forced-look head in all three games (−1.7, −4.3, −2.0; every interval spanning zero) and above the other-agent head by margins that also span zero. The reading registered for that outcome — that the second agent's places add nothing to a forced look — was written before the arm ran, and it is the reading this paper withdraws.

It is withdrawn because a matched size cannot distinguish “the second agent's groups are dead weight” from “forced-look experience saturates at half its size”, and the discriminator registered to separate them lands on the other side. A head on the mixture's own 132 forced-look groups alone solves 32.6%. So the forced arm does not saturate: 290 forced groups beat 132 by +9.0, +8.3 and +8.7 (interval excluding zero in every game; lower bounds 3.3, 2.7, 3.0 against the 3.67 bar). And the second agent's groups contribute: adding its 158 groups to the same 132 forced groups moves 32.6% to 38.6% (+7.3, +4.0, +6.7; interval excluding zero in two of three). At a matched 290 the second agent's experience looked worthless because it was displacing forced-look groups, not because it carried nothing.

The strict status of the three new pairs is the status this paper already gives the forced-look-versus-other-agent pair: none clears both bars, so each is a lean and not a ranking. The point estimates order the six heads forced 41.2 > mixture 38.6 > other-agent 35.4 > forced-half 32.6 > self 26.8 > base 22.0, and no adjacent step in that order is established: three are inside the session's instability, the self-versus-base step clears neither bar (§3), and forced-half versus self was not a registered pairing and has no interval. No ordering is claimed here.

One sentence the arm buys, stated as consistent-with rather than established, is a magnitude for the paper's variable: 132 forced-look groups are indistinguishable from 290 of the second agent's in all three games (−2.3, −1.0, −5.3). That is consistent with a forced look being worth about twice a second agent's experience per decision group in this system; it is not a measured ratio, because an interval that contains zero difference cannot license one, and it would be the first thing to test elsewhere.

6. What did not hold

- **The cheap proxy measures the wrong variable.** Held-out pick accuracy at fitting time ordered the three heads of equal size other-agent 57.9% > forced-look 54.8% > mixture 53.8%; in the loop, by problems solved and by in-game hit rate, the order is forced-look > mixture > other-agent, exactly reversed. With the fourth head the failure splits cleanly: across sizes (290 against 132 groups) the probe ranks the half-size head lowest, correctly (45.5%). The held-out probe tracks how much experience a head was fitted on and not where it came from — the one variable this paper is about. This is the sixth probe non-conversion in the series by the experimenter's count (VIII §7.6, VIII-b §13c and §4 of that paper, and the heads here), recorded in the register rather than derivable from one artefact. It is not used as a diagnostic of acquisition here, and the split is a testable claim rather than a retirement.
- **“Above the other agent” did not fire.** After one game of the forced-look arm read twenty points above the other-agent head, a fourth reading was registered for that direction. The third game came in at +3.3 with an interval spanning zero and the reading did not fire. The one-game flag is in the appendix as the kind of number this harness produces.
- **The first other-agent arm was contaminated and withdrawn.** Its head had been fitted on the very 300 problems its games scored and read 134 of 300 — a head recalling its answer key. The tell was proportion, not sign: an effect an order of magnitude larger than anything the instrument had produced. The games moved to a third set no head had seen, and every number above is from that set.

6a. Related work, as found by the author

The author is outside the research community and made no systematic survey; what follows is a first pass over about twenty-five abstracts on 2026-09-27, ranked by closeness to this paper's claim, and it is offered so that a reader who knows closer work can say so. None of the papers below holds the intelligence fixed, carries experience in a read-only channel outside the prompt under verifier labels, and varies only where the experience was acquired; that dissociation is the claim of this paper and was not found elsewhere. The instrument, however, is not new, and the two closest objects are named first.

The instrument. OLIVIA (arXiv 2605.11169) models a frozen LLM’s action selection as a contextual linear bandit over candidate actions with the frozen last-layer hidden state as context; nothing in the model is updated — the learned quantities are per-action ridge statistics kept outside it, updated online from a per-step signal that is membership of the chosen action in the task’s reference solution. It is the closest published object to VIII’s head; against it, VIII’s head reads one state at a decision point, is fitted offline from a verifier’s per-candidate region label before any head exists, and is measured on a quality–compute frontier rather than on task performance. Multi-Head Latent Control (arXiv 2607.14277) trains two heads on a frozen model’s latent traces — a capability head emitting a score on which proceed-or-defer is thresholded, from the final layer, and a resolution head choosing among clarification, tool use, abstention and direct answering, from a selected middle layer — with labels from an LLM judge against reference answers, offline; it reports the quality–cost trade-off of the resulting multi-model system. Its input is a variable-length trace over generated tokens with the prompt excluded, which is what buys its early hand-off from partial generations; this series’ head reads one state at one decision. Mid-layer reading and an offline judge label are therefore not distinctive of this series, and are not claimed as such. ExpGraph (arXiv 2605.30712) keeps a frozen executor and improves it by retrieving from an external experience graph whose nodes are summaries of the executor’s own completed trajectories, maintained by exponential moving averages of outcomes and ranked by a fixed formula of relevance and estimated utility; its one learned object is a small policy that sets two scalars of the retrieval (breadth and the relevance-versus-utility weight), trained on the paired with-and-without-retrieval score. It is the closest published work to “selecting experience for a frozen agent”, carried as memory and retrieval over task text and graph statistics rather than as a head over activations; and because every node it can retrieve was produced by the executor’s own policy, it sits inside this paper’s own-experience arm and has no analogue of the second-agent or forced-look arms. HSRM (arXiv 2608.30841) trains a small encoder over a frozen generator’s hidden states at step boundaries from self-generated trajectories with outcome labels, to rank candidate solutions in mathematics; it is the same shape as this series’ head used as a scorer. SEAM (arXiv 2602.02556) stores experience in a plug-in’s parameters and guides a frozen executor, trained by reinforcement learning on the executor’s rollouts. ReBalance (Li, Tu et al., ICLR 2026, arXiv 2603.12372) fits a training-free steering vector over a frozen reasoning model’s deep hidden states and gates it by the model’s stepwise confidence, for reasoning economy; its result with a sign-changing controller is evidence that VIII’s failed steering arm — one fixed direction, fixed strength — failed on the carrier’s form and not on the hidden state’s content, which is what this paper’s head assumes.

The plateau. A 2026 study of memory-based self-improving agents (arXiv 2608.00017) finds that without ground truth an agent re-uses its own most confident mistakes, an “echo gap” in the scoring signal; its agents learn without updating any weights, so it is a frozen-model result like this one, but its stated cause — the absence of ground truth forcing an LLM-assessed reward — does not obtain here, where labels come from a verifier, and VIII-b’s plateau is of a different kind: the states repeat, not the errors. Rise-and-collapse under self-training (arXiv 2606.21090, whose stated cause is over-optimisation on a fixed distribution) and non-compounding across iterations of experience learning (arXiv 2606.04703 — “progressive capability collapse rather than compounding improvement”, the weight-update analogue of VIII-b’s central negative result, that a second generation of the agent’s own experience is largely a repeat of the first) are the weight-update relatives; stagnation attributed to fixed task difficulty (arXiv 2609.18642) is the training-schedule relative. The paragraph groups these by symptom; at least two share a cause.

The lever. A 2026 paper decouples exploration from exploitation by having lightweight “scouts” — small MLPs, not language models — explore an environment before a language model is trained on their trajectories by supervised fine-tuning and then reinforcement learning (arXiv 2601.21754); it is the trained-model relative of this paper’s second agent, with an acquisition agent of a different kind where ours is the same frozen model under another agent’s policy, and because use is implemented by changing the model it cannot hold the consumer fixed while the source varies, which is the separation this paper makes. DORA (arXiv 2604.17244) is a training-free, inference-time rule for more diverse actions, measured by task success rather than by the value of the experience acquired; it is a natural acquisition rule for the follow-up this paper opens. ExpeL (arXiv 2308.10144) and its successors carry experience in the prompt as retrieved trajectories and distilled insights, a prompting method that leaves the model’s weights untouched; they are the baseline class a reader will ask this series to be compared against, and the generality programme names that comparison. Across the six papers above whose experience, states or labels are sourced — OLIVIA, ExpGraph, Multi-Head Latent Control, HSRM, SEAM and ExpeL — every one draws them from the model’s own behaviour, and none varies where the experience came from; the scout paper does introduce an outside explorer and changes the model’s weights to use what it finds. A fitted readout over frozen activations is therefore the field’s standard instrument; what is this paper’s is using it to vary the experience’s source with the model and the consuming mechanism both held fixed.

7. What this says, and what it does not

Learning is usually discussed as one thing. This series has been separating three: the intelligence, which is frozen throughout and never changed; the ability to *use* verified experience, which VIII showed and which every head here exercises through the same read-only form; and the ability to *acquire* experience worth using, which is what the acting policy determines by where it stands when a state is recorded. VIII-b found the plateau; this paper finds that, in this system, the plateau was in acquisition, and that changing where the agent stands before it decides — by borrowing another agent’s habit or by imposing one — restores gains of the size VIII first reported, with the intelligence held fixed.

The technical form is that the value of experience is conditional on the state distribution induced by the acquisition policy. The intuitive form is that where you learn from matters, not just how much you collect. Both dimensions are visible in the same table: quantity still matters (290 forced-look groups beat 132), and so does origin (two acquisition policies produced experience of different value per group, and each contributed). So not all experience of one form is equivalent, and the question this opens, and does not pursue, is the one the series began with, asked one level up: ET-I asked how to spend an inference budget; this paper ends by asking how to spend an experience-acquisition budget — given a fixed number of groups, which states should an agent seek in order to acquire the most useful experience? Experience acquisition is an optimisation dimension in its own right, separable from the intelligence and from the mechanism that consumes experience. Two mechanisms are consistent with the forced-look result and are not separated here: the forced look may give the head better coverage of the states the game later visits, or it may produce states in which the true region is more linearly readable, so the head learns cleaner directions. The two point to different follow-ups — seek more states, or seek more readable ones — and separating them is a registered question for VIII-d, not a claim of this paper. That is the finding of record; the second agent, the forced look and the percentages are the path that led to it.

The claim is scoped to what was measured: one task family (debugging), one model pair, one head form, 300 problems per game. “Can arise” is the verb because a mechanism has been demonstrated in this experimental system, not a law of frozen agents. The two later arms ran a day after the bracket they are scored against, and the only evidence that session drift is small is the base’s own within-day stability (66/67/65). The loop harness recorded no provenance of its own invocation at the time; the source set of the forced-look states was established from the data (candidate region names $\subseteq$ each task’s own regions: 300 of 300 against the fitting set, 15 of 300 against the scoring set) rather than from memory; the harness’s invocation record (resolved arguments, task-set and head hashes, script sha) is written and tested and is applied after the last running chain is off the file — it is not yet in the artefacts of this paper.

Reproducibility

Every game file, head, task set and score in this paper is named by sha256 in `experience/results/viic_LOCK.json` (arms of §3), `viic_LOCK_v2.json` (the forced-look and mixture arms) and `viic_LOCK_v3.json` (the half-size arm), with the git sha of the script that produced each artefact and its exact invocation. A second reader verified the first lock — 108 of 108 hashes, 16 of 16 statistics from its table, 12 of 12 green counts from the files — and found the one discrepancy that moved a recorded verdict (§6, the bar). The second and third locks (the forced-look, mixture and half-size arms) were read by the results reader from the staged files: every game file hashed against its lock entry (162 of 162, 189 of 189; zero missing, zero mismatched) and every green count recomputed from the raw episodes (39 of 39). The head files themselves are named by hash in the locks; they were not re-hashed by the second reader at the time of the lock reads, and have since been hashed against the first lock (2 of 2 exact), bringing that lock’s verifiable-from-staged total to 110 of 110; every episode those heads produced verifies. Withdrawn runs are named by hash in the first lock (36 entries); their bytes were not part of the staged verification set and rest on the experimenter’s box, and a withdrawn entry can carry the same filename as a live one, differing only by hash, so a file is identified by its hash and never by its name. The registrations, in the order they were written, are `docs/et8b-loop-gates.md §16–§16c`.

On 2026-09-27 the from-scratch script `experience/reproduce_8c.sh` completed one full run on a second machine: both task sets regenerated from their seeds, states re-collected, the forced-look head re-fitted from the run’s own states, three games per arm. Base 66, 70, 60 of 300 against the lock’s 66, 67, 65 (means 0.22 points apart); forced-look 126, 121, 116 against 122, 126, 123 (0.89 apart); the effect re-derives at +18.55 points against +19.22. Both arms are within the 3.67-point bar of §6, and the reproduced base arm’s own three-game range (10 problems) is an independent measurement that a spread of that size belongs to the loop. The stage-6 output is at `experience/results/viic_repro/2026-09-27_stage6_output.txt`; the self-experience arm is skipped by the script because its head is a VIII-b artefact not published here.

References

- Alsakka, L. (2026). *Efficient Thinking I: Measuring What Capability Costs*. This series.
- Alsakka, L. (2026). *Efficient Thinking VIII: Experience Priors*. This series.

- Alsakka, L. (2026). *Efficient Thinking VIII-b: Does Experience Accumulate?* This series.
- Asadolahi, M. et al. (2026). *Memory Reward Inflation in Self-Improving LLM Agents*. arXiv:2608.00017.
- Chen, J. et al. (2026). *Rethinking Continual Experience Internalization for Self-Evolving LLM Agents*. arXiv:2606.04703.
- Feng, T. et al. (2026). *ExpGraph: Model-Agnostic Experience Learning with Graph-Structured Memory for LLM Agents*. arXiv:2605.30712.
- Ghasemabadi, A. et al. (2026). *Multi-Head Latent Control: A Unified Interface for LLM Agent Decision Making*. arXiv:2607.14277.
- Gurjar, P. et al. (2026). *DORA Explorer: Improving the Exploration Ability of LLMs Without Training*. arXiv:2604.17244.
- Li, X. et al. (2026). *Beyond Experience Retrieval: Learning to Generate Utility-Optimized Structured Experience for Frozen LLMs*. arXiv:2602.02556.
- Li, X. et al. (2026). *HSRM: Hidden-State Reward Models for Test-Time Verification*. arXiv:2608.30841.
- Li, Y. et al. (2026). *Efficient Reasoning with Balanced Thinking*. arXiv:2603.12372. ICLR 2026.
- Lin, J. (2026). *Self-Improvement Can Self-Regress: The Rise-and-Collapse Failure Mode of LLM Self-Training*. arXiv:2606.21090.
- Wang, H. et al. (2026). *Language-based Trial and Error Falls Behind in the Era of Experience*. arXiv:2601.21754.
- Yu, S. et al. (2026). *OLIVIA: Online Learning via Inference-time Action Adaptation for Decision Making in LLM ReAct Agents*. arXiv:2605.11169.
- Yu, Y. et al. (2026). *STRETCH the Boundaries: A Unified Self-Taught Framework for Progressive LLM Evolution*. arXiv:2609.18642.
- Zhao, A. et al. (2023). *ExpeL: LLM Agents Are Experiential Learners*. arXiv:2308.10144.

Appendix A. How this was found

1. **09-23.** VIII-b’s saturation diagnosis; the question of escape posed as “a different agent, same family — the fastest route”.
2. **09-24 07:1xZ.** §16 registered: agent B collects on seed 73’s 300; the readings fixed before collection.
3. **09-24.** First collection stopped at 33 episodes: the harness prompt carried a stale region menu from an earlier task family; B obeyed it, the 7B ignored it. Corrected; a fresh 7B base run showed the correction moved vocabulary, not difficulty (22.7% vs 20.7%).
4. **09-24.** Reading (i) fired on the registered 25% bar; a same-agent churn control then showed two runs of the same agent differ by 4 problems (6.7%) — the bar was tightened to a difference against churn before any number travelled.
5. **09-24 → 09-25.** The first head games: a case-fold collision between arm names made one arm silently reuse another’s files (caught by an impossible timestamp); then the other-agent head read 134 of 300 and was withdrawn as contaminated (fitted on the scoring set). Games moved to seed 89.
6. **09-25.** Arm 1 scored: the second reader recomputed the bar from its definition and found the scorer had excluded the arm under test; the interval count fell from 2 of 3 to 1 of 3.
7. **09-25.** The owner’s forced-look control registered before arm 1’s score; its charging of the forced action found and fixed before it ran.
8. **09-26.** Forced-look scored: within the bar of the other-agent head. A one-game “above” flag registered as a fourth reading; it did not fire.

9. **09-26.** Mixture scored: below forced-look in all three games; the matched-size ambiguity named and its discriminator started the same hour.
10. **09-26.** Degeneracy tested and refuted; the selectivity finding; the probe found inverted against both outcome measures.
11. **09-26 23:0xZ.** The claim of record fixed by the owner before the draft: the plateau can arise from acquisition, not from exhaustion of use.
12. **09-26 23:51Z.** The discriminator landed against the experimenter's own six-arm reading: the forced arm does not saturate and the second agent's groups contribute. The registered "adds nothing" reading withdrawn; §5 rewritten; collection frozen at lock v3.

Credits: E ran every arm and found every instrument defect above before its number travelled; R (the results reader) found the one that moved a verdict. The owner designed the forced-look control and fixed the claim.